

Course 4: Deep Learning



IMT Atlantique
Bretagne-Pays de la Loire
École Mines-Télécom

2025-10-20

Course 4: Deep Learning

Course 4: Deep Learning



IMT Atlantique
Bretagne-Pays de la Loire
École Mines-Télécom

Summary

Last session

- Unsupervised learning - discover structure from unlabeled data
- Clustering
- Decomposition
- Preprocessing and feature selection

Today's session

- Multi-Layer Perceptron
- Convolutional Neural Networks
- Transformers

Last session

- 1 Unsupervised learning - discover structure from unlabeled data
- 2 Clustering
- 3 Decomposition
- 4 Preprocessing and feature selection

Today's session

- Multi-Layer Perceptron
- Convolutional Neural Networks
- Transformers

AI
Example: Expert rules



2025-10-20

Course 4: Deep Learning

└ Global overview...

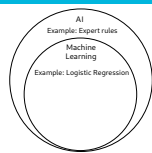
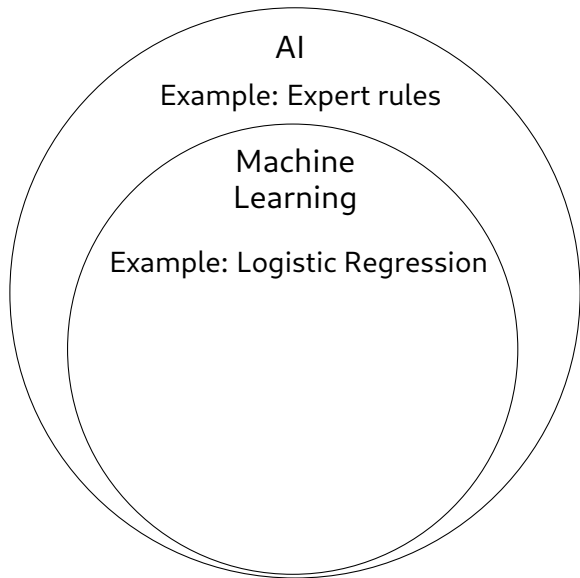
Global overview...

AI
Example: Expert rules

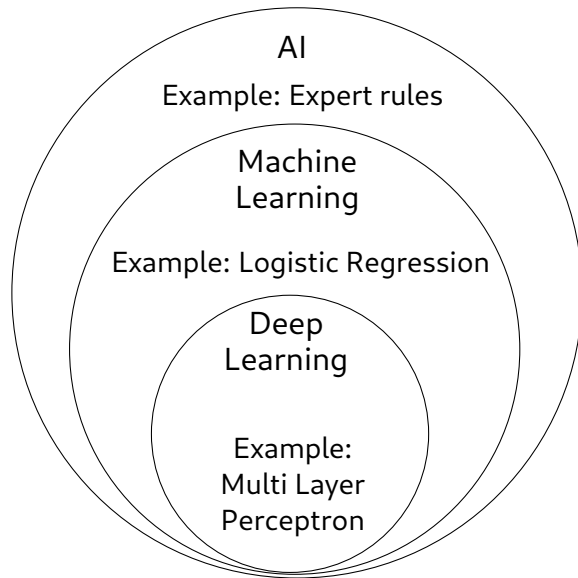


Adapted from Deep Learning, 2016 by A. Courville, I. Goodfellow and Y. Bengio

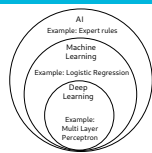
Deep Learning is a particular case of Machine Learning.




Deep Learning is a particular case of Machine Learning.




Deep Learning is a particular case of Machine Learning.





Hence, Deep Learning methods can be **supervised, unsupervised, ...**

└ Global overview...



Hence, Deep Learning methods can be supervised, unsupervised, ...

Deep Learning in a nutshell (1/3)

Tentative definition

Using *deep* Artificial Neural Networks.
Optimized on *large* datasets.

We generally talk about “Neural Networks” instead of “Artificial Neural Networks”, but the latter is the most accurate terminology!

The strength of Deep Learning lies in using a lot (*a lot*) of data.

2025-10-20

Course 4: Deep Learning

└ Deep Learning in a nutshell (1/3)

Deep Learning in a nutshell (1/3)

Tentative definition

Using *deep* Artificial Neural Networks.
Optimized on *large* datasets.

We generally talk about “Neural Networks” instead of “Artificial Neural Networks”, but the latter is the most accurate terminology!

The strength of Deep Learning lies in using a lot (*a lot*) of data.

Deep Learning in a nutshell (1/3)

Tentative definition

Using *deep* Artificial Neural Networks.
Optimized on *large* datasets.

We generally talk about “Neural Networks” instead of “Artificial Neural Networks”, but the latter is the most accurate terminology!

The strength of Deep Learning lies in using a lot (a **lot**) of data.

└ Deep Learning in a nutshell (1/3)

Tentative definition


Using deep Artificial Neural Networks.
Optimized on large datasets.

We generally talk about "Neural Networks" instead of "Artificial Neural Networks", but the latter is the most accurate terminology!

The strength of Deep Learning lies in using a lot (a lot) of data.

Deep Learning in a nutshell (2/3)

A major breakthrough in image classification:




Source: Kang, D. Y., Duong, H. P., & Park, J. C. (2020). Application of deep learning in dentistry and implantology. *Journal of implantology and applied sciences*, 24(3), 148-181.

Details for the human evaluation: Russakovsky, Dieg et al.. ImageNet Large Scale Visual Recognition Challenge, <https://arxiv.org/pdf/1409.0575.pdf>

Course 4: Deep Learning


Deep Learning in a nutshell (2/3)

A major breakthrough in image classification:



The landscape of Machine Learning changed in 2012: Deep Neural Networks, a technique used in a minority of cases until then, suddenly won the Image Classification contest on ImageNet, a standard image classification dataset. From this point, Deep Neural Networks became mainstream, and the performance of Deep Neural Network models skyrocketed.


Deep Learning in a nutshell (3/3)



Inspired from <https://www.softwaretestinghelp.com/data-mining-vs-machine-learning-vs-ai/>

2025-10-20 Course 4: Deep Learning

Deep Learning in a nutshell (3/3)



- 1 Multi-layer Perceptron
- 2 Convolutional Neural Networks (CNN)
- 3 Transformers
- 4 State of Deep Learning today

2025-10-20

Course 4: Deep Learning
└ Multi-layer Perceptron
└ Outline

Outline

■ Multi-layer Perceptron

■ Convolutional Neural Networks (CNN)


■ Transformers

■ State of Deep Learning today

Perceptron

1943, first known implementation in 1957

Perceptron is a nonlinear operation in which weights W are trainable.



Source: By Mat the w at English Wikipedia, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=23766733>


2025-10-20

Course 4: Deep Learning
└ Multi-layer Perceptron
└ Perceptron

Perceptron

1943, first known implementation in 1957

Perceptron is a nonlinear operation in which weights W are trainable.



Source: By Mat the w at English Wikipedia, CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=23766733>

The Perceptron is a matrix multiplication, followed by a nonlinear function $f()$. It is important to say that the Perceptron is old!! Hence, AI is not a brand new thing, but an old research domain.

Perceptron

1943, first known implementation in 1957

Perceptron is a nonlinear operation in which weights W are trainable.




Figure: The arrows represent the weights W .


Source: <https://davidstutz.de/illustrating-convolutional-neural-networks-in-latex-with-tikz/>

2025-10-20

Course 4: Deep Learning

└ Multi-layer Perceptron

└ Perceptron



The Perceptron is a matrix multiplication, followed by a nonlinear function $f()$. It is important to say that the Perceptron is old!! Hence, AI is not a brand new thing, but an old research domain.

Optimizing the weights

Loss

- Prediction: $y = f\left(\sum_{d=0}^D x_d W_d\right)$
- Ground truth: \hat{y}
- Loss (one example): $\mathcal{L}(x, W, \hat{y}) = d(y, \hat{y})$
(ex: $d(y, \hat{y}) = \|y - \hat{y}\|_2^2$)
- Loss (i examples): $J(W) = \sum_i \mathcal{L}(x^{(i)}, W, \hat{y}^{(i)})$

Gradient descent

- Compute the gradient: $\frac{\partial J(W)}{\partial W}$ (high dimensional derivative)
- Update weights: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$

2025-10-20

Course 4: Deep Learning

└ Multi-layer Perceptron

└ Optimizing the weights

Optimizing the weights

Loss

- Prediction: $y = f\left(\sum_{d=0}^D x_d W_d\right)$
- Ground truth: \hat{y}
- Loss (one example): $\mathcal{L}(x, W, \hat{y}) = d(y, \hat{y})$
(ex: $d(y, \hat{y}) = \|y - \hat{y}\|_2^2$)
- Loss (i examples): $J(W) = \sum_i \mathcal{L}(x^{(i)}, W, \hat{y}^{(i)})$

Gradient descent


- Compute the gradient: $\frac{\partial J(W)}{\partial W}$ (high dimensional derivative)
- Update weights: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$

The Perceptron is a matrix multiplication, followed by a nonlinear function $f()$. It is important to say that the Perceptron is old!! Hence, AI is not a brand new thing, but an old research domain.

Intuition behind the gradient descent

Update is given as: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$

- $\partial J(W)$ gives the direction
- η gives the size of the step



Adapted from <https://tex.stackexchange.com/questions/561921/replicating-a-plot-using-tikz>

2025-10-20


Course 4: Deep Learning
└ Multi-layer Perceptron
└ Gradient descent

Gradient descent

Intuition behind the gradient descent

Update is given as: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$

- $\partial J(W)$ gives the direction
- η gives the size of the step



Adapted from <https://tex.stackexchange.com/questions/561921/replicating-a-plot-using-tikz>

The gradient follows the increase of the error function, hence the inverse of the gradient follows its decrease. Parameter η , the learning rate, gives the size of the step to take at each iteration. If too small, the model will slowly converge. If too large, the model can be unstable and never reach the optimal solution. In practice, setting the learning rate is not easy.


Gradient descent

Intuition behind the gradient descent

Update is given as: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$

- $\partial J(W)$ gives the direction
- η gives the size of the step

Small step:



2025-10-20

Course 4: Deep Learning

└ Multi-layer Perceptron

└ Gradient descent


Gradient descent

Intuition behind the gradient descent

Update is given as: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$

- $\partial J(W)$ gives the direction
- η gives the size of the step

Small step:




The gradient follows the increase of the error function, hence the inverse of the gradient follows its decrease. Parameter η , the learning rate, gives the size of the step to take at each iteration. If too small, the model will slowly converge. If too large, the model can be unstable and never reach the optimal solution. In practice, setting the learning rate is not easy.

Intuition behind the gradient descent

Update is given as: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$

- $\partial J(W)$ gives the direction
- η gives the size of the step

Large step:



2025-10-20

Course 4: Deep Learning

└ Multi-layer Perceptron

└ Gradient descent

Gradient descent

Intuition behind the gradient descent

Update is given as: $W \leftarrow W - \eta \frac{\partial J(W)}{\partial W}$

■ $\partial J(W)$ gives the direction

■ η gives the size of the step

Large step:




The gradient follows the increase of the error function, hence the inverse of the gradient follows its decrease. Parameter η , the learning rate, gives the size of the step to take at each iteration. If too small, the model will slowly converge. If too large, the model can be unstable and never reach the optimal solution. In practice, setting the learning rate is not easy.

Multi-Layer Perceptron

Multi-Layer Perceptron (= *fully-connected* network)

- Stacking Perceptrons.
- The **deep** term comes from this stacking
- **Prediction:** $y = f(W^{(L)} \dots f(W^{(2)} f(W^{(1)} x)))$




Source: <https://davidstutz.de/illustrating-convolutional-neural-networks-in-latex-with-tikz/>

2025-10-20

Course 4: Deep Learning

Multi-layer Perceptron

Multi-Layer Perceptron



The bias are removed from the equations for simplicity, but say that they exist orally.

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

2025-10-20

Course 4: Deep Learning

- └ Multi-layer Perceptron
 - └ Neural Networks

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

2025-10-20

Course 4: Deep Learning

- └ Multi-layer Perceptron
 - └ Neural Networks

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the activation function.
- $\mathbf{W}^{(l)}$ are called weights.
- $\mathbf{b}^{(l)}$ is called the bias.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

2025-10-20

Course 4: Deep Learning

- └ Multi-layer Perceptron
 - └ Neural Networks

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

2025-10-20

Course 4: Deep Learning

- └ Multi-layer Perceptron
 - └ Neural Networks

Neural Networks

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

2025-10-20

Course 4: Deep Learning

- └ Multi-layer Perceptron
 - └ Neural Networks

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

Definitions

$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.

Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

2025-10-20

Course 4: Deep Learning

- └ Multi-layer Perceptron
 - └ Neural Networks

Definitions


$$\mathbf{y}^{(l+1)} = f(\mathbf{W}^{(l)}\mathbf{y}^{(l)} + \mathbf{b}^{(l)}) = \phi^{(l)}(\mathbf{y}^{(l)})$$

- Each building block $\mathbf{y}^{(l+1)}$ is called a **layer**.
- One element i of a layer ($y_i^{(l)}$) is called a **neuron** (both as input and output).
- The nonlinear function f is called the **activation function**.
- $\mathbf{W}^{(l)}$ are called **weights**.
- $\mathbf{b}^{(l)}$ is called the **bias**.


Note: while each layer can have a different activation function f , it is standard that each layer uses the same.

Why is it called Neural Network?

“Neurons” may be seen as **loosely** inspired from the human brain.



Attribution: Egm4313.s12 at English Wikipedia,
[https://en.wikipedia.org/wiki/Neural_network_\(machine_learning\)#/media/File:Neuron3.png](https://en.wikipedia.org/wiki/Neural_network_(machine_learning)#/media/File:Neuron3.png)




Detail of one neuron in an Artificial Neural Network.
* Late note: I forgot the bias in the graphic.

This is an **analogy**, artificial neural networks are **not** following the human brain (in general).

2025-10-20


Course 4: Deep Learning └ Multi-layer Perceptron

└ Why is it called Neural Network?




Why is it called Neural Network?

“Neurons” may be seen as **loosely** inspired from the human brain.



Attribution: Egm4313.s12 at English Wikipedia,
[https://en.wikipedia.org/wiki/Neural_network_\(machine_learning\)#/media/File:Neuron3.png](https://en.wikipedia.org/wiki/Neural_network_(machine_learning)#/media/File:Neuron3.png)

This is an **analogy**, artificial neural networks are **not** following the human brain (in general).




Detail of one neuron in an Artificial Neural Network.
* Late note: I forgot the bias in the graphic.

2025-10-20

Course 4: Deep Learning └ Multi-layer Perceptron


└ Why is it called Neural Network?



Gradient descent for deep neural networks

Backpropagation

- Gradient descent for all layers (chain rule).
- Simplified equation: $\frac{\partial J(W)}{\partial W} = \frac{\partial J(W)}{\partial W^{(L)}} \frac{\partial W^{(L)}}{\partial W^{(L-1)}} \frac{\partial W^{(L-1)}}{\partial W^{(L-2)}} \dots \frac{\partial W^{(2)}}{\partial W^{(1)}}$
- The error **backpropagates** through the network (reverse path)
- Computationally efficient, but finds a local minimum (at best)




Source: <http://introtodeeplearning.com/>

2025-10-20

Course 4: Deep Learning

Multi-layer Perceptron

└ Gradient descent for deep neural networks



Careful: the given equation is not correct, just conveniently simplified. For better details, refer to <https://en.wikipedia.org/wiki/Backpropagation>.

Batch

- The i examples are divided in *batches* (small excerpt)
- Allows one to train without loading the whole dataset in memory
- Accelerates the learning phase

Limits of Multi-Layer Perceptrons

- Computationally heavy for large inputs
- Large number of parameters: prone to overfitting
- No notion of structure in the input: everything is a vector

Limits of Multi-Layer Perceptrons


- Computationally heavy for large inputs
- Large number of parameters: prone to overfitting
- No notion of structure in the input: everything is a vector

- 1 Multi-layer Perceptron
- 2 Convolutional Neural Networks (CNN)
- 3 Transformers
- 4 State of Deep Learning today

Principle

- Applying a kernel to the input, on small parts of the image at a time.
- Weights of the kernel are **learned** and **shared**!
- 2D convolution was a game-changer for image processing
- Translation invariance

Convolutional layer



$$\begin{pmatrix} w_1 & w_2 & w_3 & 0 & 0 & 0 \\ 0 & w_1 & w_2 & w_3 & 0 & 0 \\ 0 & 0 & w_1 & w_2 & w_3 & 0 \\ 0 & 0 & 0 & w_1 & w_2 & w_3 \end{pmatrix}$$

2025-10-20

Course 4: Deep Learning

Convolutional Neural Networks (CNN)

Convolutional Neural Network (1/5)

Convolutional Neural Network (1/5)

Principle

- Applying a kernel to the input, on small parts of the image at a time.
- Weights of the kernel are **learned** and **shared**!
- 2D convolution was a game-changer for image processing
- Translation invariance

Convolutional layer

Convolution neural networks are the most common architecture for neural networks nowadays. They are particularly successful for image processing, were not challenged until very recently, with Vision Transformers (see Transformers in the last slides). The translation invariance is really important for image processing, as objects can be anywhere on the image. It contributed to the success of Convolutional Neural Networks.

Convolutional Neural Network (2/5)

Example of 2D convolution:

$$\begin{pmatrix} 0 & 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} * \begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 4 & 3 & 4 & 1 \\ 1 & 2 & 4 & 3 & 3 \\ 1 & 2 & 3 & 4 & 1 \\ 1 & 3 & 3 & 1 & 1 \\ 3 & 3 & 1 & 1 & 0 \end{pmatrix}$$

$I \qquad K \qquad I * K$

Source: <https://tex.stackexchange.com/questions/437007/drawing-a-convolution-with-tikz>


2025-10-20

Course 4: Deep Learning

└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (2/5)

Convolutional Neural Network (2/5)



Convolutional Neural Network (2/5)

Example of 2D convolution:

$$\begin{pmatrix} 0 & 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} * \begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 4 & 3 & 4 & 1 \\ 1 & 2 & 4 & 3 & 3 \\ 1 & 2 & 3 & 4 & 1 \\ 1 & 3 & 3 & 1 & 1 \\ 3 & 3 & 1 & 1 & 0 \end{pmatrix}$$

$I \qquad K \qquad I * K$

Source: <https://tex.stackexchange.com/questions/437007/drawing-a-convolution-with-tikz>


2025-10-20

Course 4: Deep Learning

└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (2/5)

Convolutional Neural Network (2/5)



Convolutional Neural Network (2/5)

Example of 2D convolution:

$$\begin{pmatrix} 0 & 1 & 1 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 1 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 \\ 0 & 1 & 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} * \begin{pmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{pmatrix} = \begin{pmatrix} 1 & 4 & 3 & 4 & 1 \\ 1 & 2 & 4 & 3 & 3 \\ 1 & 2 & 3 & 4 & 1 \\ 1 & 3 & 3 & 1 & 1 \\ 3 & 3 & 1 & 1 & 0 \end{pmatrix}$$

$I \qquad K \qquad I * K$


Source: <https://tex.stackexchange.com/questions/437007/drawing-a-convolution-with-tikz>

2025-10-20

Course 4: Deep Learning

└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (2/5)



Example of 2D pooling:

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 1 |
| 1 | 1 | 1 | 1 |
| 2 | 3 | 1 | 6 |
| 8 | 1 | 4 | 5 |

| | |
|--|--|
| | |
| | |

| | |
|---|--|
| 2 | |
| | |

maxpool, kernel 2, stride 2

2025-10-20

Course 4: Deep Learning

└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (3/5)

Example of 2D pooling:

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 1 |
| 1 | 1 | 1 | 1 |
| 2 | 3 | 1 | 6 |
| 8 | 1 | 4 | 5 |

| | |
|--|--|
| | |
| | |

| | |
|---|--|
| 2 | |
| | |

maxpool, kernel 2, stride 2

Example of 2D pooling:

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 1 |
| 1 | 1 | 1 | 1 |
| 2 | 3 | 1 | 6 |
| 8 | 1 | 4 | 5 |

| | |
|--|--|
| | |
| | |

| | |
|---|---|
| 2 | 3 |
| | |

maxpool, kernel 2, stride 2

2025-10-20

Course 4: Deep Learning

└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (3/5)

Example of 2D pooling:

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 1 |
| 1 | 1 | 1 | 1 |
| 2 | 3 | 1 | 6 |
| 8 | 1 | 4 | 5 |

| | |
|--|--|
| | |
| | |

| | |
|---|---|
| 2 | 3 |
| | |

maxpool, kernel 2, stride 2

Example of 2D pooling:

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 1 |
| 1 | 1 | 1 | 1 |
| 2 | 3 | 1 | 6 |
| 8 | 1 | 4 | 5 |

| | |
|--|--|
| | |
| | |

| | |
|---|---|
| 2 | 3 |
| 8 | |

maxpool, kernel 2, stride 2

2025-10-20

Course 4: Deep Learning

└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (3/5)

Example of 2D pooling:

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 1 |
| 1 | 1 | 1 | 1 |
| 2 | 3 | 1 | 6 |
| 8 | 1 | 4 | 5 |

| | |
|--|--|
| | |
| | |

| | |
|---|---|
| 2 | 3 |
| 8 | |

maxpool, kernel 2, stride 2

Example of 2D pooling:

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 1 |
| 1 | 1 | 1 | 1 |
| 2 | 3 | 1 | 6 |
| 8 | 1 | 4 | 5 |



| | |
|---|---|
| 2 | 3 |
| 8 | 6 |

maxpool, kernel 2, stride 2

2025-10-20

Course 4: Deep Learning

└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (3/5)

Example of 2D pooling:

| | | | |
|---|---|---|---|
| 1 | 2 | 3 | 1 |
| 1 | 1 | 1 | 1 |
| 2 | 3 | 1 | 6 |
| 8 | 1 | 4 | 5 |




| | |
|---|---|
| 2 | 3 |
| 8 | 6 |

maxpool, kernel 2, stride 2

Convolutional Neural Network (4/5)

And repeat...

- Convolutional neural network: mainly *Convolution + Pooling*.
- ...But many other components may be added! (batch norm, dropout, skip connections, concatenation, ...)




Source: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

2025-10-20

Course 4: Deep Learning


└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (4/5)



Why convolutions?

- Kernels capture important information in images
- The kernels become more and more complex with the depth of the network




Source: <https://ujjwalkarn.me/2016/08/11/intuitive-explanation-convnets/>

2025-10-20

Course 4: Deep Learning


└ Convolutional Neural Networks (CNN)

└ Convolutional Neural Network (5/5)



Convolution are able to catch simple shapes (lines, edges) which turn into complex shapes in the subsequent layers.

What happened in 2012?



A combination of...

- Convolutional neural networks
- A very large dataset (ImageNet)
- Clever tricks (ex: data augmentation, *i.e.* altering image during training, very standard in Deep Learning)
- The use of GPUs for computation

2025-10-20

Course 4: Deep Learning

└ Convolutional Neural Networks (CNN)

└ What happened in 2012?

What happened in 2012?



A combination of...

- Convolutional neural networks
- A very large dataset (ImageNet)
- Clever tricks (ex: data augmentation, *i.e.* altering image during training, very standard in Deep Learning)
- The use of GPUs for computation

- 1 Multi-layer Perceptron
- 2 Convolutional Neural Networks (CNN)
- 3 Transformers**
- 4 State of Deep Learning today

Transformers

2025-10-20 Course 4: Deep Learning
└ Transformers
└ Transformers

Standard architecture nowadays

- No convolution
- Based on attention: what should be important for context?
- Used for text, image, audio, ...

Standard architecture nowadays

- No convolution
- Based on *attention*: what should be important for context?
- Used for text, image, audio, ...

Standard architecture nowadays

- No convolution
- Based on *attention*: what should be important for context?
- Used for text, image, audio, ...

Will be presented in the next class (foundation models)!

Standard architecture nowadays

- No convolution
- Based on *attention*: what should be important for context?
- Used for text, image, audio, ...

Will be presented in the next class (foundation models)!

- 1 Multi-layer Perceptron
- 2 Convolutional Neural Networks (CNN)
- 3 Transformers
- 4 State of Deep Learning today

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ Outline

Outline

▣ Multi-layer Perceptron

▣ Convolutional Neural Networks (CNN)

▣ Transformers

▣ State of Deep Learning today

State-of-the-Art nowadays

Image classification

- Image classification for a single dataset is (almost) solved
- Challenges of adapting models to unseen datasets
- Challenges when data is scarce
- Specific domains with few variability or complex classification are still challenging (ex: medical imaging)

Large Language Models

- Large Language Models caught everyone's attention (ChatGPT)
- Challenges of reducing their resources (data/power)
- May hallucinate: lack of robustness

Many other domains

Multimodal models (DALL-E, ...), Audio, Games, Video, ...

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ State-of-the-Art nowadays

State-of-the-Art nowadays

Image classification

- Image classification for a single dataset is (almost) solved
- Challenges of adapting models to unseen datasets
- Challenges when data is scarce
- Specific domains with few variability or complex classification are still challenging (ex: medical imaging)

Large Language Models

- Large Language Models caught everyone's attention (ChatGPT)
- Challenges of reducing their resources (data/power)
- May hallucinate: lack of robustness

Many other domains

Multimodal models (DALL-E, ...), Audio, Games, Video, ...

State-of-the-Art nowadays

Image classification

- Image classification for a single dataset is (almost) solved
- Challenges of adapting models to unseen datasets
- Challenges when data is scarce
- Specific domains with few variability or complex classification are still challenging (ex: medical imaging)

Large Language Models

- Large Language Models caught everyone's attention (ChatGPT)
- Challenges of reducing their resources (data/power)
- May hallucinate: lack of robustness

Many other domains

Multimodal models (DALL-E, ...), Audio, Games, Video, ...

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ State-of-the-Art nowadays

State-of-the-Art nowadays

Image classification

- Image classification for a single dataset is (almost) solved
- Challenges of adapting models to unseen datasets
- Challenges when data is scarce
- Specific domains with few variability or complex classification are still challenging (ex: medical imaging)

Large Language Models

- Large Language Models caught everyone's attention (ChatGPT)
- Challenges of reducing their resources (data/power)
- May hallucinate: lack of robustness

Many other domains

Multimodal models (DALL-E, ...), Audio, Games, Video, ...

State-of-the-Art nowadays

Image classification

- Image classification for a single dataset is (almost) solved
- Challenges of adapting models to unseen datasets
- Challenges when data is scarce
- Specific domains with few variability or complex classification are still challenging (ex: medical imaging)

Large Language Models

- Large Language Models caught everyone's attention (ChatGPT)
- Challenges of reducing their resources (data/power)
- May hallucinate: lack of robustness

Many other domains

Multimodal models (DALL-E, ...), Audio, Games, Video, ...

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ State-of-the-Art nowadays

State-of-the-Art nowadays

Image classification

- Image classification for a single dataset is (almost) solved
- Challenges of adapting models to unseen datasets
- Challenges when data is scarce
- Specific domains with few variability or complex classification are still challenging (ex: medical imaging)

Large Language Models

- Large Language Models caught everyone's attention (ChatGPT)
- Challenges of reducing their resources (data/power)
- May hallucinate: lack of robustness

Many other domains

Multimodal models (DALL-E, ...), Audio, Games, Video, ...

Focus on Large Language Models

Many models

- GPT (Open-AI)
- LLaMA (Meta)
- Gemini (Google)
- Mistral 8x7B (MistralAI)
- Many others... And more to come!

Masked Language Modeling

*How are **you** doing today? → How are ... doing today?*

- The network learns to reconstruct masked words
- No supervision!
- Allows to leverage immense datasets (ex: GPT-3 was learned on an **Internet scale** dataset)

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ Focus on Large Language Models

Focus on Large Language Models

Many models

- GPT (Open-AI)
- LLaMA (Meta)
- Gemini (Google)
- Mistral 8x7B (MistralAI)
- Many others... And more to come!

Masked Language Modeling

How are you doing today? → How are ... doing today?

- The network learns to reconstruct masked words
- No supervision!
- Allows to leverage immense datasets (ex: GPT-3 was learned on an Internet scale dataset)

Focus on Large Language Models

Many models

- GPT (Open-AI)
- LLaMA (Meta)
- Gemini (Google)
- Mistral 8x7B (MistralAI)
- Many others... And more to come!

Masked Language Modeling

How are **you** doing today? → How are ... doing today?

- The network learns to reconstruct masked words
- No supervision!
- Allows to leverage immense datasets (ex: GPT-3 was learned on an **Internet scale** dataset)

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ Focus on Large Language Models

Focus on Large Language Models

Many models

- GPT (Open-AI)
- LLaMA (Meta)
- Gemini (Google)
- Mistral 8x7B (MistralAI)
- Many others... And more to come!

Masked Language Modeling

How are **you** doing today? → How are ... doing today?

- The network learns to reconstruct masked words
- No supervision!
- Allows to leverage immense datasets (ex: GPT-3 was learned on an Internet scale dataset)

Focus on Large Language Models

Many models

- GPT (Open-AI)
- LLaMA (Meta)
- Gemini (Google)
- Mistral 8x7B (MistralAI)
- Many others... And more to come!

Masked Language Modeling

How are **you** doing today? → How are ... doing today?

- The network learns to reconstruct masked words
- No supervision!
- Allows to leverage immense datasets (ex: GPT-3 was learned on an **Internet scale** dataset)

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ Focus on Large Language Models

Focus on Large Language Models

Many models

- GPT (Open-AI)
- LLaMA (Meta)
- Gemini (Google)
- Mistral 8x7B (MistralAI)
- Many others... And more to come!

Masked Language Modeling

How are **you** doing today? → How are ... doing today?


- The network learns to reconstruct masked words
- No supervision!
- Allows to leverage immense datasets (ex: GPT-3 was learned on an **Internet scale** dataset)

Large Language Models are greedy

Model sizes

- AlexNet (2012): 62 Million parameters
- GPT-3 (2020): 175 Billion parameters


Image source:
<https://informationisbeautiful.net/visualizations/the-rise-of-generative-ai-large-language-models/>



2025-10-20

Course 4: Deep Learning
└ State of Deep Learning today

└ Large Language Models are greedy



Conclusion

- Deep Learning algorithms: **powerful** without feature extraction
- They require **a lot** of data to be trained
- The architecture plays an important role

Common criticisms

- Hard to interpret
- Reproduce biases from data
- May require massive amounts of energetic consumption

Going further

- Details and maths behind IA: <https://youtu.be/aircAruvnKk>
- Ethics and reflexions (french): Science4all & M.Phi

<https://youtu.be/HbFadt0xs4k> | https://youtu.be/_XJsAQsTOBo

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ Deep Learning

Conclusion

- Deep Learning algorithms: **powerful** without feature extraction
- They require **a lot** of data to be trained
- The architecture plays an important role

Common criticisms

- Hard to interpret
- Reproduce biases from data
- May require massive amounts of energetic consumption

Going further

- Details and maths behind IA: <https://youtu.be/aircAruvnKk>
- Ethics and reflexions (french): Science4all & M.Phi

Conclusion

- Deep Learning algorithms: **powerful** without feature extraction
- They require **a lot** of data to be trained
- The architecture plays an important role

Common criticisms

- Hard to interpret
- Reproduce biases from data
- May require massive amounts of energetic consumption

Going further

- Details and maths behind IA: <https://youtu.be/aircAruvnKk>
- Ethics and reflexions (french): Science4all & M.Phi

<https://youtu.be/HbFadt0xs4k> | https://youtu.be/_XJsAQsTOBo

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ Deep Learning

Deep Learning

Conclusion

- Deep Learning algorithms: **powerful** without feature extraction
- They require **a lot** of data to be trained
- The architecture plays an important role

Common criticisms

- Hard to interpret
- Reproduce biases from data
- May require massive amounts of energetic consumption

Going further

- Details and maths behind IA: <https://youtu.be/aircAruvnKk>
- Ethics and reflexions (french): Science4all & M.Phi

<https://youtu.be/HbFadt0xs4k> | https://youtu.be/_XJsAQsTOBo

Conclusion

- Deep Learning algorithms: **powerful** without feature extraction
- They require **a lot** of data to be trained
- The architecture plays an important role

Common criticisms

- Hard to interpret
- Reproduce biases from data
- May require massive amounts of energetic consumption

Going further

- Details and maths behind IA: <https://youtu.be/aircAruvnKk>
- Ethics and reflexions (french): Science4all & M.Phi

<https://youtu.be/HbFadt0xs4k> | https://youtu.be/_XJsAQsT0Bo

2025-10-20

Course 4: Deep Learning

└ State of Deep Learning today

└ Deep Learning

Deep Learning

Conclusion

- Deep Learning algorithms: **powerful** without feature extraction
- They require **a lot** of data to be trained
- The architecture plays an important role

Common criticisms

- Hard to interpret
- Reproduce biases from data
- May require massive amounts of energetic consumption

Going further

- Details and maths behind IA: <https://youtu.be/aircAruvnKk>
- Ethics and reflexions (french): Science4all & M.Phi

<https://youtu.be/HbFadt0xs4k> | https://youtu.be/_XJsAQsT0Bo

Lab

- Lab Pytorch: manipulating the basics of PyTorch
- Lab Modality: try a first baseline on your modality